ASSIGNMENT 7 – FINAL PROJECT

REAL TIME ANALYTICS

Kumari Yachana

CU23MSD0013A

MSc. Data Science

School of Mathematics and Natural Sciences

Project Report: Real-Time Weather Monitoring and Visualization System

Introduction

The Real-Time Weather Monitoring and Visualization System is designed to provide live updates on weather conditions using an end-to-end pipeline. The project fetches weather data from a public API, processes it using Apache Kafka, stores it in a PostgreSQL database, and visualizes the data using Apache Superset. This system facilitates monitoring, analyzing, and presenting weather metrics, enabling stakeholders to make informed decisions.

Project Plan

Objectives:

1. Build a pipeline to fetch live weather data.
2. Process and stream data using Apache Kafka.
3. Store the transformed data in PostgreSQL.
4. Design a real-time, auto-refreshing dashboard using Apache Superset.
5. Enable seamless data consumption and visualization.

Steps:

1. **Data Collection** : Fetch weather data from the OpenWeatherMap API.
2. **Data Processing** : Stream data via Kafka Producer and Consumer.
3. **Data Storage** : Load data into PostgreSQL for persistence.
4. **Visualization** : Create an interactive dashboard to present the data.
5. **System Validation** : Test for accuracy, latency, and usability.

System Architecture Diagram :

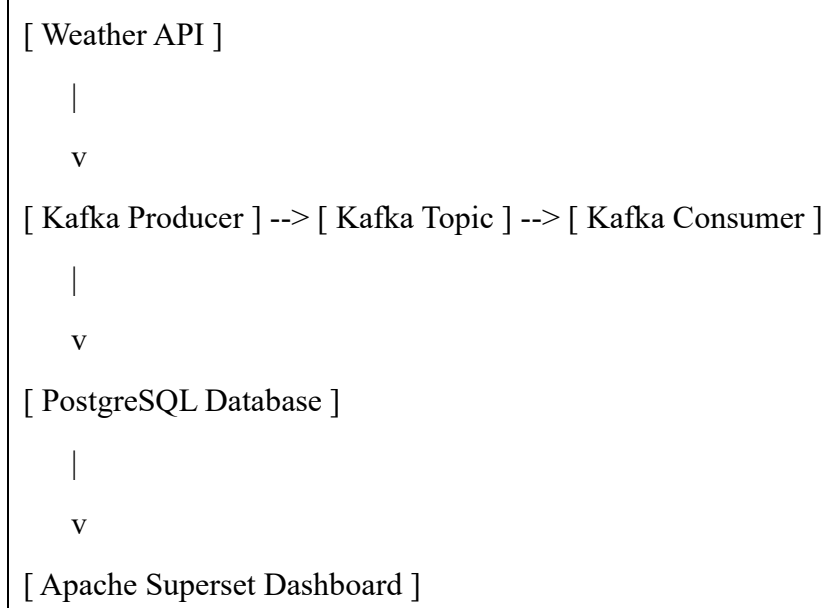

Metrics Used in Weather Monitoring

1. **Temperature :**
   - Unit: Degrees Celsius (or Fahrenheit).
   - Metric: Real-time temperature of a city.
2. **Humidity :**
   - Unit: Percentage (%).
   - Metric: Represents atmospheric moisture content.
3. **Weather Condition :**
   - Descriptions: Clear, Cloudy, Rainy, etc.
   - Metric: Categorical classification of weather.
4. **Timestamp :**
   - Unit: Date-Time.
   - Metric: Represents the exact time of data capture.

Building the End-to-End System

1. **Data Collection**

- Fetched data from the OpenWeatherMap API using an API key.
- Query parameters: City, API key, and units.

2. **Data Transformation**

- **Kafka Producer** :
- Fetches and sends JSON data to a Kafka Topic.
- Example payload:

json

```
{
  "city": "London",
  "temperature": 12.5,
  "humidity": 80,
  "weather_description": "Clear",
  "timestamp": "2025-01-16T10:00:00Z"
}
```

```
kafka@hp-HP-Laptop-15s-eq2xxx:/home/hp/kafka$ /home/hp/kafka/bin/kafka-topics.sh --create --topic weather-data --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
Created topic weather-data.
kafka@hp-HP-Laptop-15s-eq2xxx:/home/hp/kafka$
```

```
hp@hp-HP-Laptop-15s-eq2xxx: $ /home/hp/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092
test-topic
weather-data
hp@hp-HP-Laptop-15s-eq2xxx: $
```

Python code for producer:

Open

weatherProducer.py
~/

```python
from kafka import KafkaProducer
import requests
import json
import time

API_KEY = "cd1fae386adff5a987f22a4fae910791"
API_URL = "http://api.openweathermap.org/data/2.5/weather"
LOCATION = "London"  # Change to your desired location
KAFKA_BROKER = "localhost:9092"
KAFKA_TOPIC = "weather-data"

def fetch_weather_data():
    params = {
        'q': LOCATION,
        'appid': API_KEY,
        'units': 'metric'
    }
    response = requests.get(API_URL, params=params)
    response.raise_for_status()
    return response.json()

def main():
    producer = KafkaProducer(
        bootstrap_servers=KAFKA_BROKER,
        value_serializer=lambda v: json.dumps(v).encode('utf-8')
    )
    while True:
        try:
            weather_data = fetch_weather_data()
            producer.send(KAFKA_TOPIC, value=weather_data)
            print(f"Sent data: {weather_data}")
        except Exception as e:
            print(f"Error fetching or sending data: {e}")
        time.sleep(10)  # Adjust the interval as needed

if __name__ == "__main__":
    main()
```

Running the code :

- **Kafka Consumer :**

- Consumes data from the Kafka Topic.

- Transforms data into a format compatible with PostgreSQL.

Python code for consumer :

Open

weatherConsumer.py
~/

```python
from kafka import KafkaConsumer
import psycopg2
import json

KAFKA_BROKER = "localhost:9092"
KAFKA_TOPIC = "weather-data"
POSTGRES_CONN = {
    'dbname': 'weatherdb',
    'user': 'yachana',
    'password': 'yachana6',
    'host': 'localhost',
    'port': 5432
}

def create_table():
    conn = psycopg2.connect(**POSTGRES_CONN)
    cursor = conn.cursor()
    cursor.execute("""
        CREATE TABLE IF NOT EXISTS weather_data (
            id SERIAL PRIMARY KEY,
            city TEXT,
            temperature REAL,
            humidity INT,
            weather_description TEXT,
            timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP
        );
    """)
    conn.commit()
    cursor.close()
    conn.close()

def save_to_postgres(data):
    conn = psycopg2.connect(**POSTGRES_CONN)
    cursor = conn.cursor()
    cursor.execute("""
        INSERT INTO weather_data (city, temperature, humidity, weather_description)
        VALUES (%s, %s, %s, %s);
    """, (data['name'], data['main']['temp'], data['main']['humidity'], data['weather'][0]['description']))
    conn.commit()
    cursor.close()
    conn.close()

def main():
    create_table()
    consumer = KafkaConsumer(
        KAFKA_TOPIC,
        bootstrap_servers=KAFKA_BROKER,
        value_deserializer=lambda v: json.loads(v.decode('utf-8'))
    )
    for message in consumer:
        try:
            save_to_postgres(message.value)
            print(f"Saved data: {message.value}")
        except Exception as e:
```

Running the code :

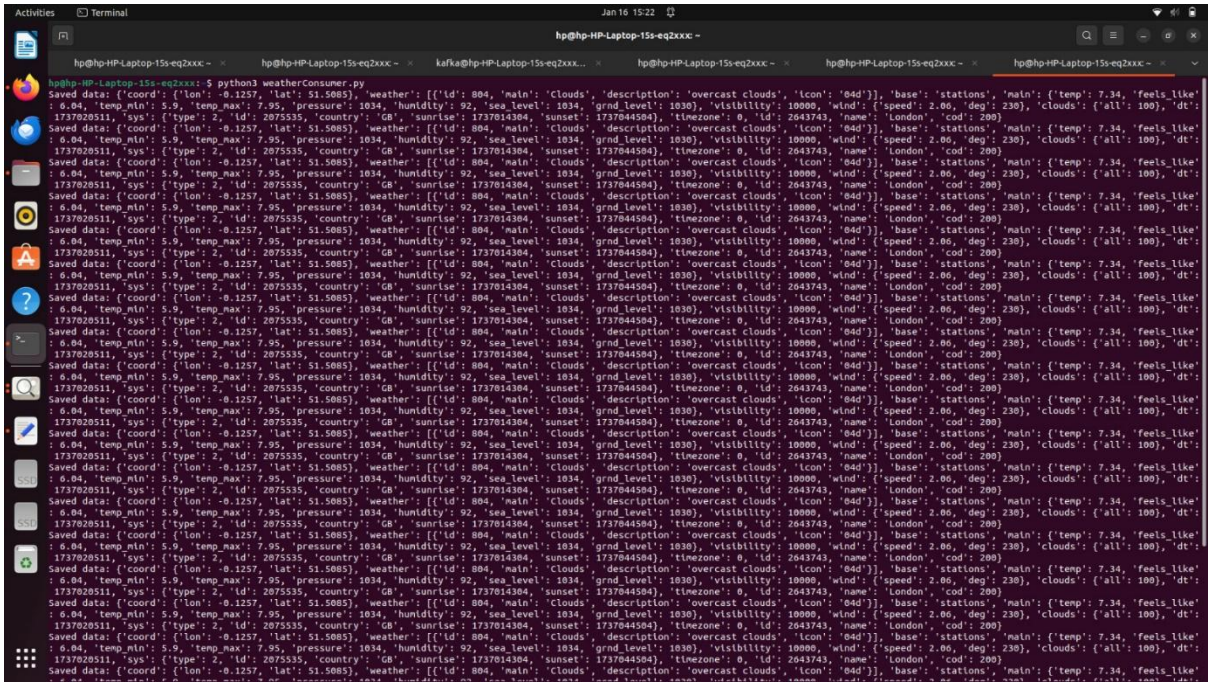

3. **Data Storage**

- PostgreSQL :

- Schema Design:

sql

```
CREATE TABLE weather_data (
    id SERIAL PRIMARY KEY,
    city TEXT,
    temperature REAL,
    humidity INT,
    weather_description TEXT,
    timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

- Insert transformed data into the database.

Created database “weatherdb” and table “weather_data” in it :

The data generated is inserted into the table :

```
You are now connected to database "weatherdb" as user "yachana".
weatherdb=# \dt
            List of relations
 Schema |     Name     | Type  |  Owner
--------+--------------+-------+---------
 public | weather_data | table | yachana
(1 row)

weatherdb=# SELECT * FROM weather_data;
 id |  city  | temperature | humidity | weather_description |         timestamp
----+--------+-------------+----------+---------------------+----------------------------
  1 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:18:19.371339
  2 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:18:29.476603
  3 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:18:39.660496
  4 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:18:49.829669
  5 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:18:59.95678
  6 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:19:10.095944
  7 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:19:20.26384
  8 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:19:31.80934
  9 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:19:41.920251
 10 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:19:52.032157
 11 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:20:02.156331
 12 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:20:17.257291
 13 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:20:27.36825
 14 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:20:38.433014
 15 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:20:53.582783
 16 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:21:03.750989
 17 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:21:13.855165
 18 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:21:24.031659
 19 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:21:34.392223
 20 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:21:44.521342
 21 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:21:55.087051
 22 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:05.193019
 23 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:15.33514
 24 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:25.472405
 25 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:35.573982
 26 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:45.744347
 27 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:22:55.848021
 28 | London |        7.46 |       92 | overcast clouds     | 2025-01-16 15:23:05.988042
 29 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:23:16.117391
 30 | London |        7.34 |       92 | overcast clouds     | 2025-01-16 15:23:26.22403
```

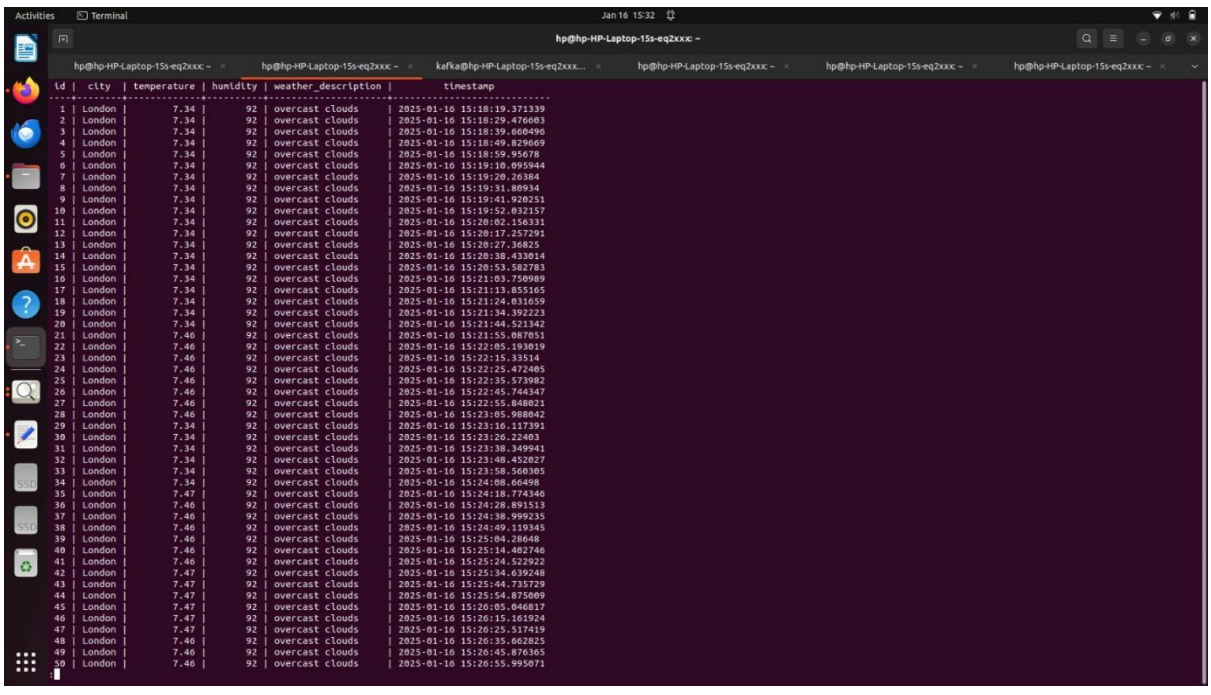

4. **Data Subscription :**

- Kafka enables real-time streaming, ensuring new weather updates are consistently pushed into PostgreSQL.

5. **Data Presentation** :-

-**Apache Superset** :

- Creating visualizations (e.g., time-series charts, heatmaps, and gauge charts).
- Configuring auto-refresh to update charts every 5-30 seconds.
- Example Dashboards:
  - Line chart showing temperature trends.
  - Heatmap visualizing humidity across cities.
  - Gauge chart displaying live temperature.

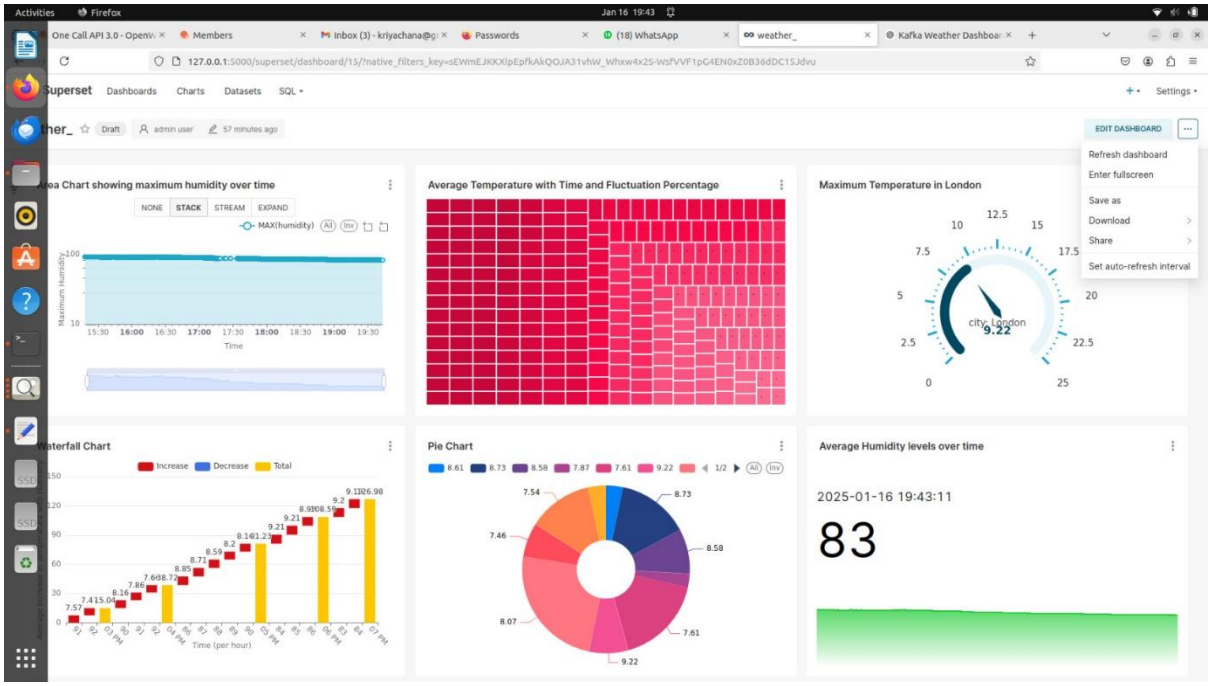

Auto refreshing the charts in 10 seconds :

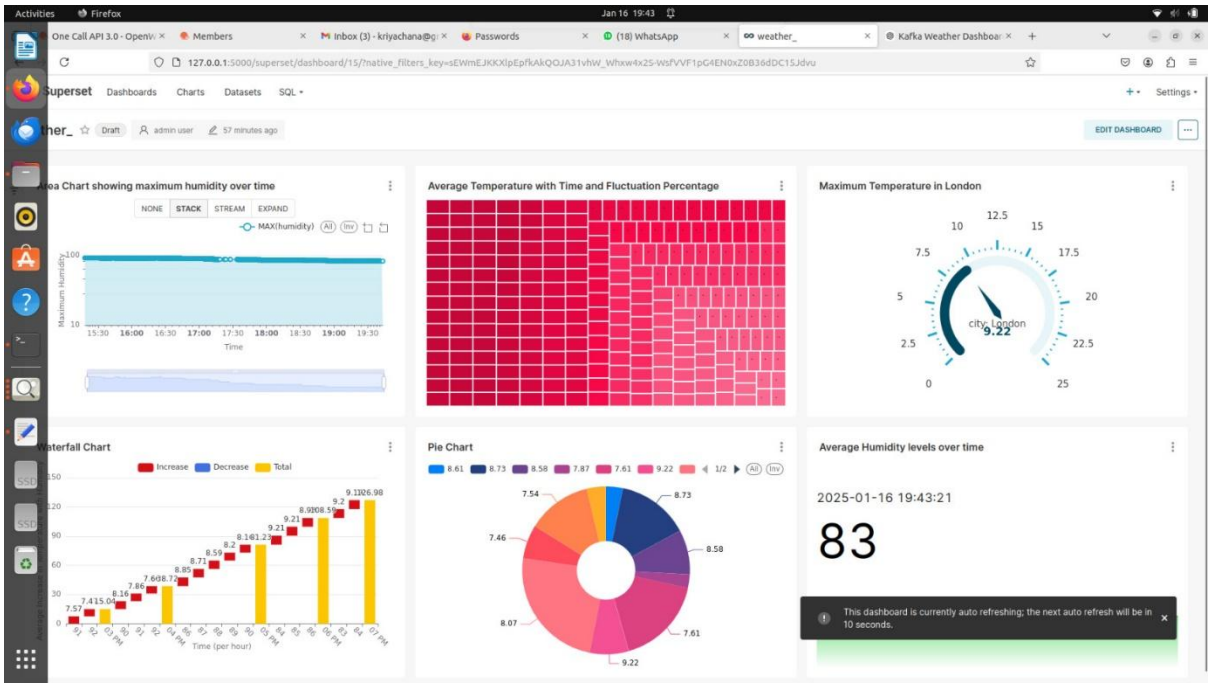

Benefits and Conclusion

This project provides an efficient system for real-time weather monitoring and visualization. By leveraging Kafka for streaming, PostgreSQL for persistence, and Superset for presentation, stakeholders can access up-to-date weather insights. The system is scalable, reliable, and user-friendly, making it an ideal solution for applications requiring live data analysis and decision-making.

If further enhancements or new metrics are needed, the modular design of this system allows for easy upgrades and scaling.